

STEMKit: A client-side toolkit for computational chemistry and scientific data analysis

Olanrewaju M. Daramola^{a,1,*}

^a*Independent Researcher*

Abstract

STEMKit is a suite of 21 browser-based tools for computational chemistry and scientific data analysis, built on `@stemkit/core`, a JavaScript library of parsing and numerical routines carrying no dependency on the document object model. All computation executes inside the user's own browser: no data is transmitted, no account is required, and nothing is installed, so the tools remain usable for unpublished or confidential material. The identical modules also execute under Node.js, so an analysis explored interactively can be captured as a version-pinned script and placed under continuous integration. Sixteen domain modules are covered by 1075 tests validated against SciPy, NumPy, and physical invariants rather than against the implementation itself.

Keywords: Computational chemistry, Molecular dynamics, Client-side computation, GROMACS, Reproducibility, Data privacy

Required Metadata

Current code version

1. Motivation and significance

The software landscape of computational chemistry is well served at its extremes. Simulation engines such as GROMACS [1] and LAMMPS [2] are mature, extensively validated, and actively maintained; analysis frameworks

*Corresponding author.

Email address: lanrelangmuir@gmail.com (Olanrewaju M. Daramola)

¹ORCID: 0009-0006-3327-2047

Nr.	Code metadata description	Metadata
C1	Current code version	v0.1.0
C2	Permanent link to code/repository used for this code version	https://github.com/LD-Shell/stemkit (release tag v0.1.0)
C3	Legal code license	MIT License
C4	Code versioning system used	git
C5	Software code languages, tools and services used	JavaScript (ECMAScript 2020 modules), HTML5, CSS3; Node.js; Jest; Tailwind CSS; GitHub Actions
C6	Compilation requirements, operating environments and dependencies	Browser tools: any current browser supporting ECMAScript modules and the <code>FileReader</code> API; no installation and no build step. Library: Node.js ≥ 18 ; <code>npm install</code> installs the development dependencies (Jest, Tailwind CSS) only. There are no runtime dependencies — <code>jStat</code> , <code>Papa Parse</code> , <code>regression.js</code> and <code>bibtex-parse-js</code> are vendored in the repository
C7	If available, link to developer documentation/manual	https://github.com/LD-Shell/stemkit#readme ; per-module API documentation in <code>src/-core/README.md</code> ; hosted tools at https://stemkit.net
C8	Support email for questions	lanrelangmuir@gmail.com (issue tracker: https://github.com/LD-Shell/stemkit/issues)

Table 1: Code metadata (mandatory)

such as MDAnalysis [4] provide comprehensive programmatic access to trajectory data. Between these poles lies a substantial volume of routine work that neither category serves comfortably.

Consider the interval between completing a simulation and reporting it. A researcher must inspect the RMSD to confirm equilibration, extract a radius

of gyration, convert a structure between the nanometre convention of GRO-MACS and the Å convention of the Protein Data Bank, compare replicate measurements statistically, digitise a figure from a paper for comparison, abbreviate journal titles in a reference list, and compose a scheduler script requesting an appropriate resource allocation. None of these operations is intellectually demanding. All of them are error-prone when performed by hand, and all of them are performed repeatedly.

1.1. The friction problem

Three classes of solution exist, and each imposes a characteristic cost.

Local scripting. A Python environment with MDAnalysis, NumPy and SciPy addresses every task listed above and is the correct instrument for sustained analytical work. Its cost is environmental. Managed institutional workstations frequently prohibit package installation; conda environments drift; and a script written six months ago may no longer execute against the current interpreter. For a researcher whose primary expertise lies in chemistry rather than software engineering, this cost is not negligible.

Web services. Server-side tools eliminate installation entirely but require data transmission. For unpublished results, material subject to a confidentiality agreement, or work under embargo, this is frequently prohibited outright. Even where permitted, the transfer creates a record whose retention policy the researcher does not control.

Desktop applications. Graphical applications avoid both problems but introduce a third: they are difficult to script. An analysis performed by clicking cannot readily be re-executed, version-controlled, or subjected to continuous integration. In practice such an analysis is difficult to reproduce, even when every step was recorded at the time.

1.2. The client-side position

Modern browsers provide a capable computational environment. The **FileReader** API grants access to local files without transmission, and JavaScript engines execute numerical code fast enough for the data volumes involved. ECMAScript modules make it possible to organise that code properly, instead of accumulating scripts.

STEMKit is built around these constraints. No server participates in the computation, so there is nothing to which data could be transmitted and users are not asked to take a privacy policy on trust. The reproducibility

property derives from a separate architectural decision: the numerical core is extracted into modules with no dependency on the browser, so the identical code that executes behind the interface also executes under Node.js.

1.3. *Experimental setting*

The software is used in two ways, which share one implementation. Interactively, the user opens a tool page at <https://stemkit.net> (or a local clone) and selects a local file — a GROMACS `.xvg` series, a PDB or `.gro` structure, a CSV of replicate measurements — and reads the result in the page. The file is read through `FileReader`; it is never uploaded. Programmatically, the same routines are imported from `@stemkit/core` in a Node.js script, which is the path taken once an exploratory analysis is to be fixed, version-pinned and re-run. Every tool that produces a figure additionally emits a standalone matplotlib script that regenerates that figure from the original data file, so an interactive result can be reproduced offline without the browser.

Related work is cited where the corresponding functionality is described in Section 2: GROMACS [1] and LAMMPS [2] for the trajectory and job-script formats, PLUMED [3] for collective-variable output, MDAnalysis [4] as the reference framework for trajectory analysis, SciPy [5] and NumPy [6] as the validation references for the numerics, and the primary statistical literature [8, 9, 10, 11, 12, 13] for the implemented tests. Conversion factors follow CODATA 2018 [7], and the numerical treatment of tail probabilities follows standard practice for the incomplete beta and gamma functions [14].

2. Software description

2.1. *Software architecture*

The original implementation followed the pattern common to browser tooling: each of the 21 tools consisted of an HTML document paired with a single JavaScript file combining event handling, rendering and computation. This arrangement is expedient during development and progressively obstructive thereafter. Computational logic entangled with DOM manipulation cannot be tested without a browser, cannot be reused across tools, and cannot be invoked from a script.

The refactoring applied the *strangler fig* pattern: for each tool, the mathematical and parsing logic was extracted into a pure module beneath `src/-core/`, a test suite was written against the extracted module, and the original file was then reduced to DOM wiring that imports the new functions. The

user-visible interface is unchanged; the computation beneath it is now independently addressable. Figure 1 shows the resulting layering and Table 2 enumerates the modules.




Figure 1: Architecture of STEMKit. Both hosts execute the identical core modules; the only environment-specific code is the registration call that supplies the vendored libraries.

Four third-party libraries are vendored within the repository: jStat for statistical distributions, Papa Parse for delimited-text parsing, regression.js for curve fitting, and bibtex-parse-js for reference parsing. Each is distributed as a Universal Module Definition (UMD) bundle, which detects its host at run time and either assigns `module.exports` under CommonJS or attaches a global to `window` in a browser. Neither path is reachable from a plain ECMAScript module: a UMD file contains no `export` statements, so a direct import yields no bindings, while `import { createRequire } from 'module'` — the conventional Node.js workaround — constitutes a hard *resolution* failure in a browser, which occurs before any code executes and therefore cannot be intercepted by `try/catch`. Rather than branch on the execution environment inside every module, which would render the core untestable in one of its two targets, the library declares the libraries it requires and each host registers them at initialisation (Section 2.3). Modules request their dependencies lazily, at the point of use rather than at import time, so importing any module never fails merely because an unrelated library has not yet been registered.

One consequence merits explicit note, as it is a common failure mode. A

Table 2: Modules of `@stemkit/core`. Vendored dependencies are supplied by injection; modules marked *none* operate without any third-party code.

Module	Domain	Tests	Dependency
<code>bibtex</code>	Reference deduplication, sanitising	102	<code>bibtex-parse-js</code>
<code>structure</code>	Molecular geometry, PDB/GRO/XYZ	100	<i>none</i>
<code>statistics</code>	Inferential statistics	82	<code>jStat</code>
<code>curve-fitting</code>	Least-squares regression	77	<code>regression.js</code>
<code>xvg-parser</code>	GROMACS/PLUMED trajectory data	77	<i>none</i>
<code>units</code>	Physical unit conversion	73	<i>none</i>
<code>plumed</code>	PLUMED input generation	72	<i>none</i>
<code>selection</code>	Atom selection, spatial queries	67	<i>none</i>
<code>slurm</code>	HPC job-script generation	61	<i>none</i>
<code>latex</code>	Table generation and escaping	59	<i>none</i>
<code>error-bars</code>	Group summaries, error bars	57	<code>jStat</code>
<code>journals</code>	Whole-title journal abbreviation	53	<i>none</i>
<code>data-cleaning</code>	Tabular transformation	51	<code>Papa Parse</code>
<code>outliers</code>	Anomaly detection	51	<code>jStat</code>
<code>digitizer</code>	Figure digitisation	49	<i>none</i>
<code>iso4</code>	ISO 4 word-level abbreviation (LTWA)	44	<i>none</i>
Total		1075	

repository declaring `"type": "module"` instructs Node.js to parse every `.js` file beneath it as an ECMAScript module, including the vendored UMD bundles; under that interpretation the UMD factory takes its browser branch and fails. Scoping `"type": "commonjs"` to the dependency directory alone restores correct behaviour.

The test suite comprises 1075 tests across the 16 domain modules, with 93.9% statement and 97.6% line coverage of `src/core/`. Its governing principle is that numerical results are validated against *independent* references rather than against the implementation under test, since a test written from the same source as the code confirms only internal consistency. Reference values were obtained from SciPy [5] for t-tests, ANOVA, correlation, non-parametric tests, quantiles and multiple-comparison correction; from NumPy [6] for regression coefficients and descriptive statistics; from `scipy.constants` for every unit conversion factor; and from physical invariants for structural geometry — specifically the molecular weight and centre of mass of water, the orthonormality of rotation matrices, the preservation of interatomic distances under rotation, and round-trip fidelity through every supported file format. An end-to-end smoke test additionally exercises one representative path through fifteen of the sixteen modules against a real installation, detecting the class of failure that unit tests cannot: a broken aggregate export,

a misconfigured module type declaration, or a vendored bundle that fails to load.

2.2. Software functionalities

2.2.1. Trajectory and collective-variable data

GROMACS writes analysis output in the Grace format, a plain-text convention in which lines beginning with `#` are comments, lines beginning with `@` are formatting directives, and all remaining non-empty lines are whitespace-delimited numeric records. The directives carry the plot title, axis labels and per-series legends. Discarding them — as a naive parser that skips all non-numeric lines would — loses precisely the metadata that identifies which column contains which observable. The `xvg-parser` module recovers this metadata alongside the numeric matrix. Malformed records are counted and reported, not silently discarded, so a truncated trajectory is visible to the user instead of quietly shortening the analysis. Two further details of the format are handled explicitly: the Grace multi-set separator `&` is treated as a delimiter rather than parsed as data, and records containing `NaN` or `Infinity` are rejected whole, since one non-finite entry would otherwise contaminate every downstream statistic.

PLUMED [3] writes `COLVAR` files whose header lines begin with `#!` and carry field names and metadata. Although the convention differs from Grace, the comment character coincides, and the same parser therefore handles both formats without modification. This is verified by an explicit test rather than left to assumption.

2.2.2. Structure files and molecular geometry

The `structure` module parses the three coordinate formats in routine use for classical simulation: PDB and GROMACS `.gro`, both fixed-column formats in which fields are located by character position, and XYZ, which is whitespace-delimited. Fixed-column parsing is essential rather than pedantic: a residue name may legitimately be blank, and a whitespace-splitting parser would silently shift every subsequent field on such a line. Unit handling is explicit throughout — PDB and XYZ express coordinates in Å, `.gro` uses nanometres, and every parse result carries its unit — because silently mixing the two is the most direct route to a structure that is wrong by a factor of ten.

Assigning an element to each atom is a prerequisite for any mass-weighted quantity, and PDB atom names are ambiguous by construction: `CA` denotes

the α -carbon in a protein residue and calcium in an ion record. The resolution implemented here proceeds in five stages.

1. An explicit element column, where present, takes priority.
2. A leading digit is a hydrogen-count prefix (1HB, 2HG1) and is stripped before the symbol is read.
3. An unambiguous two-letter metal (FE, ZN, SE) is recognised wherever it appears, including within a host residue such as haem or selenomethionine.
4. Otherwise, a two-letter prefix whose second character is uppercase is a PDB remoteness indicator, and only the first letter is the element: CA in a protein residue is C α .
5. As a final fallback, a two-letter symbol is trusted when the residue shares its name, the convention for monatomic ions.

Stages 2 and 3 correct defects present in the original implementation (Section 2.2.5). The resolution is validated against 40 atom names drawn from real structures, covering protein backbone and side-chain atoms, nucleic acid atoms, water, monatomic ions and metal centres. On this basis the module computes the geometric centroid, the mass-weighted centre of mass, the axis-aligned bounding box, and the radius of gyration

$$R_g = \sqrt{\frac{\sum_i m_i |\mathbf{r}_i - \mathbf{R}|^2}{\sum_i m_i}}, \quad \mathbf{R} = \frac{\sum_i m_i \mathbf{r}_i}{\sum_i m_i}, \quad (1)$$

together with rigid-body rotation about an arbitrary pivot using intrinsic Z – Y – X Euler angles. Velocities present in a `.gro` file rotate with the frame but are never translated, so a round trip through rotation and export preserves the kinetic state.

2.2.3. Statistics, outliers and curve fitting

The `statistics` module implements Welch’s [8] and Student’s t-tests, the paired t-test, one-way ANOVA, Pearson correlation, the Mann–Whitney U test and the Wilcoxon signed-rank test. Welch’s test is the default for independent samples, as it does not assume equal variances and loses negligible power when they are in fact equal. Every parametric test reports an effect size and, where standard, a confidence interval. Assumptions are checked and surfaced rather than presumed: normality by the D’Agostino–Pearson K^2 omnibus test [9], and homogeneity of variance by the median-centred Brown–Forsythe variant of Levene’s test [10], chosen for its markedly greater robustness to non-normality than the original mean-centred formulation. Where

an assumption fails, a non-parametric alternative is recommended. Table 3 maps each computed quantity to its numerical method and validation reference.

Table 3: Numerical methods and validation references.

Quantity	Method	Validated against
Student / Welch t , p	Regularised incomplete beta $I_x(a, b)$	<code>scipy.stats.ttest_ind</code>
Welch–Satterthwaite ν	Closed form	SciPy, hand calculation
One-way ANOVA F , p	Sum-of-squares decomposition; I_x tail	<code>scipy.stats.f_oneway</code>
Pearson r , CI	Fisher z transform	<code>scipy.stats.pearsonr</code>
Mann–Whitney U	Tie-corrected normal approximation	<code>scipy.stats.mannwhitneyu</code>
Wilcoxon W	Signed-rank, zeros discarded	<code>scipy.stats.wilcoxon</code>
D’Agostino–Pearson K^2	Skewness and kurtosis transforms	<code>scipy.stats.normaltest</code>
Levene W	Brown–Forsythe (median-centred)	<code>scipy.stats.levene</code>
Quantiles	Linear interpolation (type 7)	<code>numpy.percentile</code>
Multiple comparisons	Holm–Bonferroni [11]	Independent implementation
Grubbs G	Inverted t critical value	<code>scipy.stats.t.ppf</code>
Modified Z -score	Median/MAD, 0.6745 scaling [12]	<code>numpy</code>
Least-squares fitting	Normal equations (regression.js)	<code>numpy.polyfit</code>

Upper-tail probabilities are computed through the complementary form of the incomplete beta and gamma functions rather than as $1 - F(x)$. The two expressions are algebraically identical; numerically they are not. Once $F(x)$ rounds to unity in double precision, the subtraction cancels completely and the result is exactly zero. The consequence is not academic. For the one-way ANOVA used in the test suite ($F = 377.545$ on 2 and 21 degrees of freedom), the naive form returns $p = 0$, whereas

$$p = I_{\frac{d_2}{d_2+d_1F}}\left(\frac{d_2}{2}, \frac{d_1}{2}\right) = 3.461\,174\,738\,291\,4 \times 10^{-17}, \quad (2)$$

agreeing with SciPy to twelve significant figures. A continuous test statistic

cannot produce $p = 0$; when that value appears in published output it indicates a numerical fault in the software, not an unusually strong effect. The same treatment is applied to the Student t , χ^2 and normal tails. Beyond $|z| \approx 8$ the vendored `erfc` implementation underflows, and a documented asymptotic expansion takes over so that a p -value is reported as 1.5×10^{-23} rather than as zero.

The `outliers` module provides Tukey fences, the Iglewicz–Hoaglin modified Z-score [12] and Grubbs’ test [13]; `curve-fitting` performs least-squares fitting with model-adequacy warnings. One inherited limitation is documented rather than fixed, and it applies to two of the three transformed models rather than to all of them. The vendored `regression.js` library fits the *power* model by regressing $\ln y$ on $\ln x$ unweighted, so the residuals minimised are those in logarithmic space and small y values carry more influence than a direct fit would give them. The *exponential* model also works in log space but weights each point by y , which is the standard correction for log-transform bias and keeps the result closer to a direct non-linear fit. The *logarithmic* model carries none of this bias: it regresses y on $\ln x$, so only the predictor is transformed and the residuals minimised are those in y , exactly as for any model linear in its parameters. Where the bias does apply it is material — for a clean doubling series the weighted exponential fit returns a growth rate of 0.690 22 where unweighted log-space least squares gives $\ln 2 = 0.693\,15$. Neither value is incorrect, since they answer different questions, but the difference is large enough to matter when a rate constant is reported. `fitCurve` therefore exposes a `linearised` flag, true for the exponential and power models only, so callers may surface the distinction; users requiring publication-grade non-linear fits are directed to Levenberg–Marquardt optimisation on untransformed data. A dataset containing points a model cannot transform is refused with a message naming the requirement — exponential needs every $y > 0$, power every $x > 0$ and $y > 0$, logarithmic only $x > 0$ — rather than having the offending points silently dropped.

2.2.4. Digitisation, unit conversion and authoring utilities

Recovering numerical values from a published figure requires a two-point affine calibration per axis. On a logarithmic axis the interpolation must be performed in logarithmic space,

$$v(p) = 10^{\log_{10} v_1 + \frac{(p - p_1)(\log_{10} v_2 - \log_{10} v_1)}{p_2 - p_1}}, \quad (3)$$

because a decade occupies a constant pixel distance there. Treating a logarithmic axis linearly is the most frequent source of error in figure digitisation,

and it produces values that are plausible rather than obviously wrong. The module additionally reports the data-space uncertainty corresponding to a one-pixel positioning error, since the precision of a digitised point is set by the resolution of the figure and not by the number of decimal places the export happens to print.

Sixty-four units across ten categories are supported, each defined by its ratio to a category base unit, so conversion reduces to $x_{\text{to}} = x_{\text{from}} \cdot f_{\text{to}}/f_{\text{from}}$. This keeps the table linear rather than quadratic in the number of units and renders every entry independently checkable against its cited source; factors are CODATA 2018 values [7], each verified against `scipy.constants` during testing. Temperature is handled separately, since the Celsius and Fahrenheit scales are affine rather than multiplicative.

Composing a scheduler script correctly requires knowledge that is engine-specific and easily misapplied, and two aspects dominate the resulting waste of allocation. First, the resource model differs by engine: GROMACS is threaded and requires one MPI rank per node with many CPUs per task, whereas LAMMPS [2] is MPI-parallel and requires many ranks per node. Emitting the wrong shape either idles most of a node or oversubscribes it. Second, and less widely appreciated, the SLURM `--mem` directive in an array job specifies memory *per task*, not per job: a request of 32 GB across a hundred-task array asks the scheduler for 3.2 TB concurrently. The generator computes the true in-flight total and reports it. Generated scripts include checkpoint-resume logic, so a queued job continues from its last checkpoint rather than silently restarting from $t = 0$ — a failure mode that consumes an entire allocation while appearing to succeed. Warnings are returned as structured objects rather than formatted strings, leaving presentation to the caller. The remaining modules support manuscript preparation: LaTeX and Markdown table generation with correct escaping, ISO-4 journal title abbreviation, and BibTeX deduplication by union-find over normalised DOIs and titles.

2.2.5. Defects exposed by extraction

Separating computation from presentation subjected the numerical code to independent verification for the first time, and three defects were identified that had propagated into reported output. First, skewness and kurtosis are defined against the *population* standard deviation,

$$G_1 = \frac{\sqrt{n(n-1)}}{n-2} \cdot \frac{1}{n} \sum_{i=1}^n \left(\frac{x_i - \bar{x}}{\sigma_{\text{pop}}} \right)^3, \quad \sigma_{\text{pop}} = \sqrt{\frac{1}{n} \sum_i (x_i - \bar{x})^2}, \quad (4)$$

whereas the original implementation used the Bessel-corrected sample standard deviation, deflating skewness by a factor of $((n - 1)/n)^{3/2}$ — approximately 15% at $n = 10$. Because both moments feed the D’Agostino–Pearson statistic, every normality p -value the tool reported was affected. Second, all upper-tail probabilities were computed by subtraction and floored at zero for strong effects. Third, haem iron was assigned the mass of fluorine (18.998 Da rather than 55.845 Da), selenomethionine selenium the mass of sulfur, and numeric-prefixed hydrogens no mass at all, so molecular weights and centres of mass computed for metalloproteins were correspondingly wrong. Each defect is corrected and covered by a regression test. Users who generated figures with earlier versions of these tools should re-check any reported skewness, normality p -value, or metalloprotein mass.

2.3. Sample code snippets analysis

Listing 1 contains the whole of the environment-specific code in the library. Each host registers the vendored bundles once, after which the identical modules run unchanged in both.

```
// Node.js: the vendored bundles are loaded through
createRequire.
import { createRequire } from 'module';
import { registerVendor } from '@stemkit/core';

const require = createRequire(import.meta.url);
registerVendor({ jStat: require('./js/dependencies/jstat.min.js') });

// Browser: the UMD <script> tags have already installed
their globals.
import { registerFromGlobals } from './src/core/index.js';
registerFromGlobals();
```

Listing 1: Registration of the vendored libraries under Node.js and in the browser.

3. Illustrative examples

The example below is the post-simulation triage that motivated the toolkit: confirm that a production run has equilibrated, characterise the final structure, test a replicate comparison and queue a continuation. Every step is available interactively as a tool page; the scripted form is shown because it is the reproducible one.

Listing 2 reads a GROMACS RMSD series. The parser returns the numeric matrix together with the title and per-series legends recovered from the @ directives, so each column is reported under its own observable name rather than by index.

```
import { parseXvg, extractColumn, columnStats } from '@stemkit/core';
import { readFileSync } from 'fs';

const { matrix, headers, title } = parseXvg(readFileSync('rmsd.xvg', 'utf8'));
// title -> "RMSD & Radius of Gyration"
// headers -> ["Time (ps)", "Backbone RMSD", "Rg"]

headers.slice(1).forEach((name, i) => {
  const s = columnStats(extractColumn(matrix, i + 1));
  console.log(`${name}: ${s.mean.toFixed(4)} +/- ${s.std.toFixed(4)} nm`);
});
```

Listing 2: Parsing a GROMACS RMSD trajectory and summarising each series.

Listing 3 takes the final coordinates and computes the mass-weighted centre of mass and radius of gyration. Because the input is a .gro file its coordinates are in nanometres, and the conversion to Å is requested explicitly rather than assumed.

```
import { parseStructure, centreOfMass, radiusOfGyration,
  convert }
  from '@stemkit/core';

const frame = parseStructure(readFileSync('final.gro', 'utf8'), 'gro');
const com = centreOfMass(frame.atoms); // nm, mass-weighted
const rg = radiusOfGyration(frame.atoms, com); // nm

console.log(`Rg = ${convert(rg, 'nm', 'angstrom').toFixed(2)} Å`);
```

Listing 3: Geometry of the final frame, with an explicit unit conversion.

Listing 4 compares two sets of replicate measurements. The returned object carries the assumption checks alongside the test result, so a violated assumption is visible at the point of use rather than discovered later; here a failed normality check redirects the comparison to Mann–Whitney.

```
import { welchTest, mannWhitneyU } from '@stemkit/core';
```

```

const r = welchTest(wildType, mutant);
// r.p, r.effectSize (Hedges' g), r.ci, r.assumptions

if (!r.assumptions.normality.passed) {
  console.warn(r.assumptions.normality.recommendation);
  console.log(mannWhitneyU(wildType, mutant));
} else {
  console.log(`Welch t(${r.df.toFixed(1)}) = ${r.t.toFixed(3)}, p = ${r.p}`);
}

```

Listing 4: A replicate comparison with the assumption checks surfaced.

Finally, Listing 5 generates the submission script for the continuation run. The engine argument selects the GROMACS resource shape — one rank per node, many CPUs per task — and the returned warnings report the true in-flight memory total for the array as well as any request the scheduler is likely to reject.

```

import { generateScript } from '@stemkit/core';

const { script, warnings } = generateScript({
  engine: 'gromacs', jobName: 'prod_md', partition: 'gpu',
  nodes: 1, gpus: 1, cpusPerTask: 16,
  walltime: '24:00:00', memory: '32G',
  modules: ['gcc/11.3', 'cuda/12.1', 'gromacs/2023.3'],
  tpr: 'md.tpr', deffnm: 'md', maxh: 23.5
});

warnings.forEach(w => console.warn(`[${w.level}] ${w.message}`));

```

Listing 5: Generating a GROMACS submission script with resource warnings.

The same four steps performed through the browser produce identical numbers, because they call the same functions; the difference is only that the browser supplies the file through `FileReader` and renders the result. Where a step produces a figure, the tool also emits a matplotlib script that regenerates it from the original data file.

4. Impact

The most direct effect of the architecture is on work that cannot currently be done at all with server-based tools. Analyses of unpublished results, of

material under a confidentiality agreement, and of data held under an institutional policy that forbids third-party upload are all outside the reach of a web service, and are frequently outside the reach of local scripting too where the workstation is managed and package installation is blocked. Because no data leaves the browser, these constraints are satisfied structurally rather than contractually: there is no server to receive the data and consequently no retention policy to evaluate. For researchers in that position the relevant comparison is not with a faster tool but with performing the calculation by hand.

For research questions already being pursued, the improvement is in reproducibility and in the reliability of the numbers. Extracting the computation into modules that run unchanged under Node.js means that an analysis first explored by clicking can be captured as a script, version-pinned and placed under continuous integration without being rewritten in another language — the step at which reproducibility is most often lost. The emitted matplotlib scripts close the same loop for figures. Independently of that, the validation strategy raises the floor: every reported statistic is checked against SciPy or NumPy rather than against the authors' expectations, and the test suite is intended to be run by users evaluating the software, not solely by its maintainer, since a library whose numerical claims cannot be checked by its users offers reproducibility only in principle.

The defects found during extraction (Section 2.2.5) carry a wider lesson, and are the most useful outcome of this work. All three had survived in software that ran without error and produced plausible numbers. None would have been caught by testing the code against expectations derived from the same source; each surfaced only when results were compared with an independent implementation. A 15% error in a reported skewness is not visible by inspection, and neither is a p -value floored at zero, which looks like a very strong result. Authors of scientific software — particularly of the small, in-house analysis code that is rarely reviewed — would be well advised to check their numerics against an independent implementation before reporting output.

In daily practice, the change is the removal of a setup step from tasks that previously carried one. Checking whether an RMSD has converged, or whether a replicate difference survives an assumption check, becomes an operation performed at the moment the question arises rather than one deferred until an environment is available. The tools are also usable in teaching and in shared computing environments where installation is not an option, since a URL is the only prerequisite.

STEMKit is released under the MIT licence, hosted at <https://stemkit>.

net, developed in the open at <https://github.com/LD-Shell/stemkit> and archived for citation. It is not used in a commercial setting and has not led to a spin-off company; the client-side design does, however, make it directly applicable to industrial research groups whose data-handling policies preclude uploading results to a third-party service.

5. Conclusions

STEMKit is a suite of 21 browser-based tools for computational chemistry built on a tested, dependency-injected JavaScript core of 16 modules. The architecture addresses three constraints simultaneously that existing tooling addresses only in pairs: it requires no installation, transmits no data and remains scriptable.

The privacy property is structural rather than contractual, because computation occurs entirely within the browser and there is no server to receive data. The reproducibility property follows from the separation of computation from presentation: the identical modules that execute behind the browser interface execute under Node.js, so an interactive analysis may be captured as a version-pinned script without rewriting. The 1075-test suite validates every numerical result against an independent reference, and the three defects that this strategy exposed — deflated standardised moments, tail probabilities floored at zero, and misassigned elements in metalloproteins — are documented, corrected and covered by regression tests. Future work will extend format coverage in the trajectory and structure modules and replace the two log-space fits with a Levenberg–Marquardt implementation on untransformed data.

Acknowledgements

I thank the maintainers of jStat, Papa Parse, regression.js, Plotly, KaTeX and 3Dmol.js, whose libraries STEMKit builds upon.

Declaration of competing interest

The author declares no competing financial or non-financial interests.

References

- [1] M. J. Abraham, T. Murtola, R. Schulz, S. Páll, J. C. Smith, B. Hess, and E. Lindahl, GROMACS: High performance molecular simulations through multi-level parallelism from laptops to supercomputers, *SoftwareX* **1–2**, 19–25 (2015). <https://doi.org/10.1016/j.softx.2015.06.001>
- [2] A. P. Thompson, H. M. Aktulga, R. Berger, D. S. Bolintineanu, W. M. Brown, P. S. Crozier, P. J. in 't Veld, A. Kohlmeyer, S. G. Moore, T. D. Nguyen, R. Shan, M. J. Stevens, J. Tranchida, C. Trott, and S. J. Plimpton, LAMMPS — a flexible simulation tool for particle-based materials modeling at the atomic, meso, and continuum scales, *Comput. Phys. Commun.* **271**, 108171 (2022). <https://doi.org/10.1016/j.cpc.2021.108171>
- [3] G. A. Tribello, M. Bonomi, D. Branduardi, C. Camilloni, and G. Bussi, PLUMED 2: New feathers for an old bird, *Comput. Phys. Commun.* **185**(2), 604–613 (2014). <https://doi.org/10.1016/j.cpc.2013.09.018>
- [4] N. Michaud-Agrawal, E. J. Denning, T. B. Woolf, and O. Beckstein, MDAnalysis: A toolkit for the analysis of molecular dynamics simulations, *J. Comput. Chem.* **32**(10), 2319–2327 (2011). <https://doi.org/10.1002/jcc.21787>
- [5] P. Virtanen, R. Gommers, T. E. Oliphant, M. Haberland, T. Reddy, D. Cournapeau, *et al.*, SciPy 1.0: fundamental algorithms for scientific computing in Python, *Nat. Methods* **17**(3), 261–272 (2020). <https://doi.org/10.1038/s41592-019-0686-2>
- [6] C. R. Harris, K. J. Millman, S. J. van der Walt, R. Gommers, P. Virtanen, D. Cournapeau, *et al.*, Array programming with NumPy, *Nature* **585**(7825), 357–362 (2020). <https://doi.org/10.1038/s41586-020-2649-2>
- [7] E. Tiesinga, P. J. Mohr, D. B. Newell, and B. N. Taylor, CODATA recommended values of the fundamental physical constants: 2018, *Rev. Mod. Phys.* **93**(2), 025010 (2021). <https://doi.org/10.1103/RevModPhys.93.025010>
- [8] B. L. Welch, The generalization of ‘Student’s’ problem when several different population variances are involved, *Biometrika* **34**(1–2), 28–35 (1947). <https://doi.org/10.1093/biomet/34.1-2.28>

- [9] R. B. D’Agostino, An omnibus test of normality for moderate and large size samples, *Biometrika* **58**(2), 341–348 (1971). <https://doi.org/10.1093/biomet/58.2.341>
- [10] M. B. Brown and A. B. Forsythe, Robust tests for the equality of variances, *J. Am. Stat. Assoc.* **69**(346), 364–367 (1974). <https://doi.org/10.1080/01621459.1974.10482955>
- [11] S. Holm, A simple sequentially rejective multiple test procedure, *Scand. J. Stat.* **6**(2), 65–70 (1979).
- [12] B. Iglewicz and D. C. Hoaglin, *How to Detect and Handle Outliers*, ASQC Quality Press, Milwaukee, WI (1993).
- [13] F. E. Grubbs, Procedures for detecting outlying observations in samples, *Technometrics* **11**(1), 1–21 (1969). <https://doi.org/10.1080/00401706.1969.10490657>
- [14] W. H. Press, S. A. Teukolsky, W. T. Vetterling, and B. P. Flannery, *Numerical Recipes: The Art of Scientific Computing*, 3rd ed., Cambridge University Press, Cambridge (2007).